

Unit V: Distributed Data Storage and Big Data

References:

- 1.** G. Coulouris, J. Dollimore and T. Kindberg; **Distributed Systems Concepts and Design,4th Edition.**
- 2.** Andrew S. Tanenbaum and Maarten van Steen; **Distributed Systems: Principles and Paradigms, 2nd Edition.**

Distributed Data Storage and Big Data

5.1 Distributed File Systems

- 5.1.1 Google File System (GFS):
Architecture, Operations, features and use cases
- 5.1.2 Hadoop Distributed File System (HDFS):
Architecture, Operations, features and use cases

5.2 NoSQL Databases:

- 5.2.1 Cassandra: Architecture, design principles,
Data Model and Query Language, features and use cases
- 5.2.2 MongoDB: Architecture, design principles,
Data Model and Query Language, features and use cases

5.3 Big Data Processing Frameworks

- 5.3.1 Hadoop: Key features and use cases
- 5.3.2 Spark : Key Features and use cases

DISTRIBUTED FILE SYSTEMS

- File system were originally developed for centralized computer systems and desktop computers.
- File system was as an operating system facility providing a convenient programming interface to disk storage.
- A **Distributed File System** (DFS) is simply a classical model of a file system distributed across multiple machines. The purpose is to promote sharing of dispersed files.
- The resources on a particular machine are **local** to itself. Resources on other machines are **remote**.
- A file system provides a service for clients. The server interface is the normal set of file operations: create, read, etc. on files.

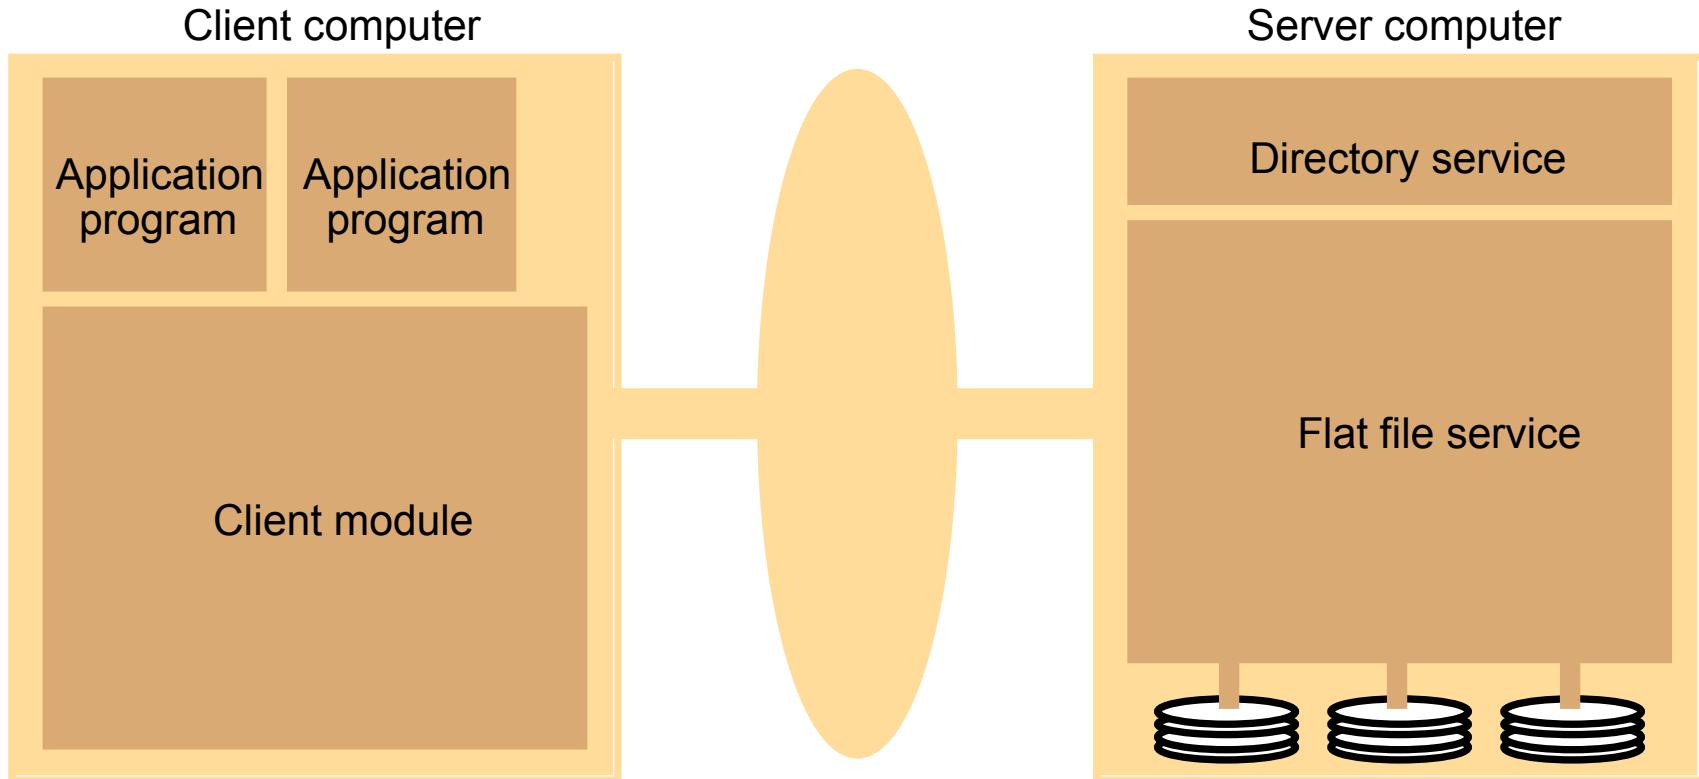
DISTRIBUTED FILE SYSTEMS

- Distributed file systems support the sharing of information in the form of files throughout the intranet.
- A distributed file system enables programs to store and access remote files exactly as they do on local ones, allowing users to access files from any computer on the intranet.
- Recent advances in higher bandwidth connectivity of switched local networks and disk organization have lead high performance and highly scalable file systems.

DISTRIBUTED FILE SYSTEMS

- Clients, servers, and storage are dispersed across machines. Configuration and implementation may vary -
 - a) Servers may run on dedicated machines, OR
 - b) Servers and clients can be on the same machines.
 - c) The OS itself can be distributed (with the file system a part of that distribution).
 - d) A distribution layer can be interposed between a conventional OS and the file system.
- Clients should view a DFS the same way they would a centralized FS; the distribution is hidden at a lower level.
- Performance is concerned with throughput and response time.

DISTRIBUTED FILE SYSTEMS



File service architecture

Distributed File system requirements

- Related requirements in distributed file systems are:
 - ❖ Transparency
 - ❖ Concurrency
 - ❖ Replication
 - ❖ Heterogeneity
 - ❖ Fault tolerance
 - ❖ Consistency
 - ❖ Security
 - ❖ Efficiency

Naming and Transparency

Naming is the mapping between logical and physical objects.

- Example: A user filename maps to <cylinder, sector>.
- In a conventional file system, it's understood where the file actually resides; the system and disk are known.
- In a **transparent** DFS, the location of a file, somewhere in the network, is hidden.
- **File replication** means multiple copies of a file; mapping returns a SET of locations for the replicas.

Location transparency -

- a) The name of a file does not reveal any hint of the file's physical storage location.
- b) File name still denotes a specific, although hidden, set of physical disk blocks.

Remote File Access

Remote File Access

CACHING

- Reduce network traffic by retaining recently accessed disk blocks in a cache, so that repeated accesses to the same information can be handled locally.
- If required data is not already cached, a copy of data is brought from the server to the user.
- Perform accesses on the cached copy.
- Files are identified with one master copy residing at the server machine,
- Copies of (parts of) the file are scattered in different caches.
- **Cache Consistency Problem** -- Keeping the cached copies consistent with the master file.
- A remote service ((RPC)) has these characteristic steps:
 - a) The client makes a request for file access.
 - b) The request is passed to the server in message format.
 - c) The server makes the file access.
 - d) Return messages bring the result back to the client.
- This is equivalent to performing a disk access for each request.

Remote File Access

Remote File Access

STATEFUL VS. STATELESS SERVICE:

Stateful: A server keeps track of information about client requests.

- It maintains what files are opened by a client; connection identifiers; server caches.
- Memory must be reclaimed when client closes file or when client dies.

Stateless: Each client request provides complete information needed by the server (i.e., filename, file offset).

- The server can maintain information on behalf of the client, but it's not required.

Remote File Access

Remote File Access

STATEFUL VS. STATELESS SERVICE:

Performance is better for stateful.

- Don't need to parse the filename each time, or "open/close" file on every request.

Fault Tolerance: A stateful server loses everything when it crashes.

- Server must poll clients in order to renew its state.
- Client crashes force the server to clean up its encached information.
- Stateless remembers nothing so it can start easily after a crash.

Remote File Access

Remote File Access

FILE REPLICATION:

- Duplicating files on multiple machines improves availability and performance.
- Placed on failure-independent machines (they won't fail together).
- The main problem is consistency - when one copy changes, how do other copies reflect that change? Often there is a tradeoff: consistency versus availability and performance.

Google File System (GFS)

Google Disk Farm



Early days...

...today



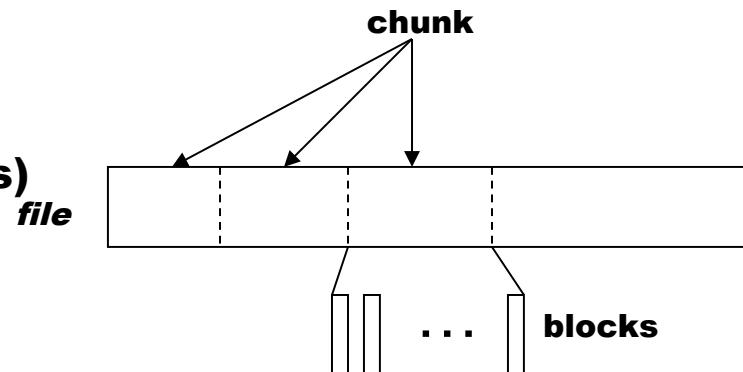
Design

■ Design factors

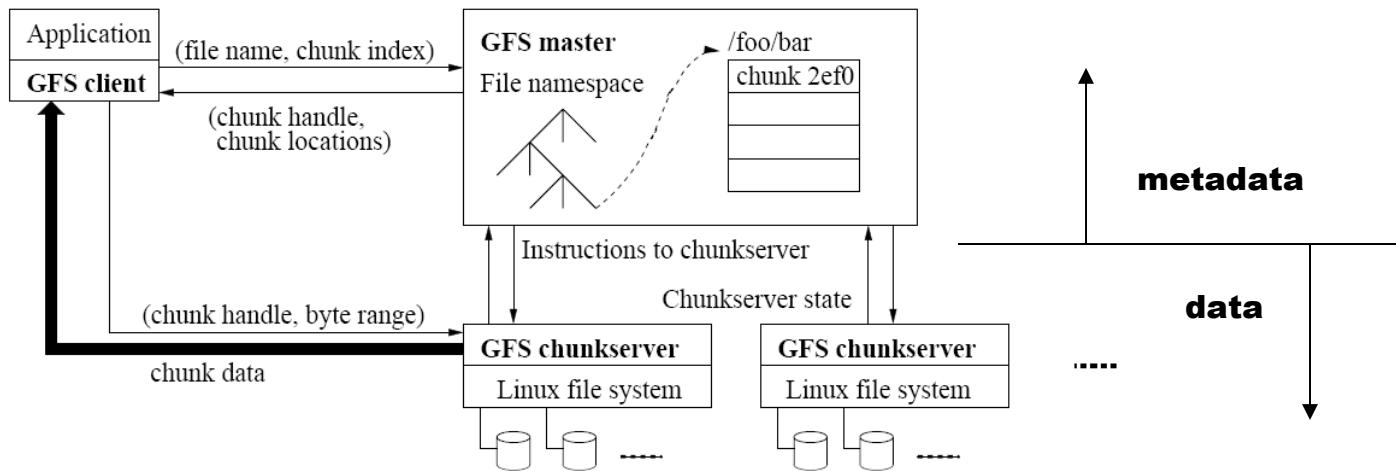
- Failures are common (built from inexpensive commodity components)
- Files
 - ❖ large (multi-GB)
 - ❖ mutation principally via appending new data
 - ❖ low-overhead atomicity essential
- Co-design applications and file system API
- Sustained bandwidth more critical than low latency

■ File structure

- **Divided into 64 MB chunks**
- **Chunk identified by 64-bit handle**
- **Chunks replicated (default 3 replicas)**
- **Chunks divided into 64KB blocks**
- **Each block has a 32-bit checksum**



Architecture



- **Master**
 - Manages namespace/metadata
 - Manages chunk creation, replication, placement
 - Performs snapshot operation to create duplicate of file or directory tree
 - Performs checkpointing and logging of changes to metadata
- **Chunkservers**
 - Stores chunk data and checksum for each block
 - On startup/failure recovery, reports chunks to master
 - Periodically reports sub-set of chunks to master (to detect no longer needed chunks)

Mutation operations

- Primary replica

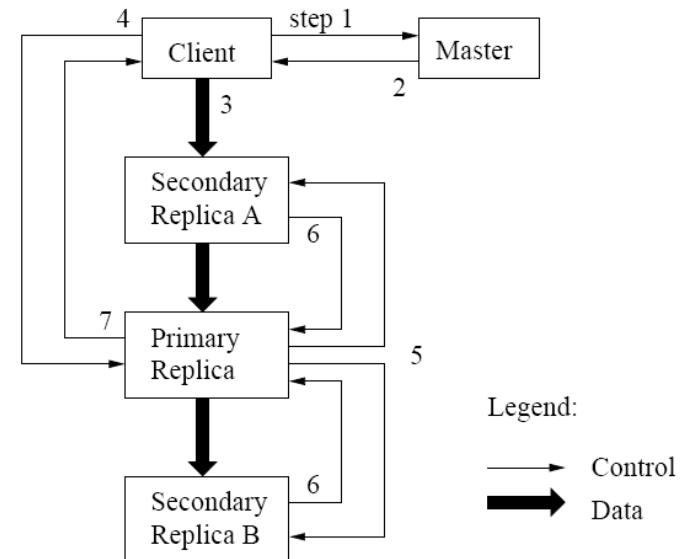
- Holds lease assigned by master (60 sec. default)
- Assigns serial order for all mutation operations performed on replicas

- Write operation

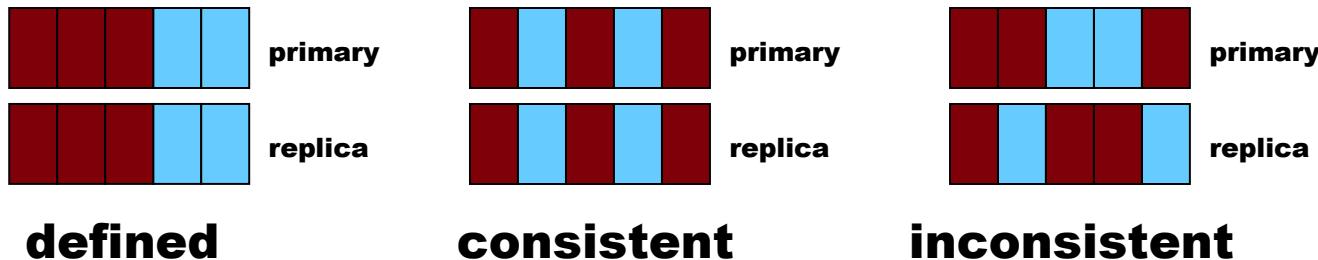
- 1-2: client obtains replica locations and identity of primary replica
- 3: client pushes data to replicas (stored in LRU buffer by chunk servers holding replicas)
- 4: client issues update request to primary
- 5: primary forwards/perform write request
- 6: primary receives replies from replica
- 7: primary replies to client

- Record append operation

- Performed atomically (one byte sequence)
- At-least-once semantics
- Append location chosen by GFS and returned to client
- Extension to step 5:
 - ❖ If record fits in current chunk: write record and tell replicas the offset
 - ❖ If record exceeds chunk: pad the chunk, reply to client to use next chunk



Consistency Guarantees



- Write
 - **Concurrent writes may be consistent but undefined**
 - **Write operations that are large or cross chunk boundaries are subdivided by client into individual writes**
 - **Concurrent writes may become interleaved**
- Record append
 - **Atomically, at-least-once semantics**
 - **Client retries failed operation**
 - **After successful retry, replicas are defined in region of append but may have intervening undefined regions**
- Application safeguards
 - **Use record append rather than write**
 - **Insert checksums in record headers to detect fragments**
 - **Insert sequence numbers to detect duplicates**

	Write	Record Append
Serial success	<i>defined</i>	<i>defined</i> interspersed with <i>inconsistent</i>
Concurrent successes	<i>consistent</i> but <i>undefined</i>	
Failure	<i>inconsistent</i>	

Metadata management

Logical structure

pathname	lock	chunk list
/home	read	Chunk4400488, ...
/save		Chunk8ffe07783, ...
/home/user/foo	write	Chunk6254ee0, ...
/home/user	read	Chunk88f703, ...

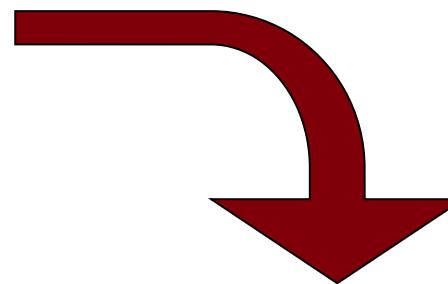
- Namespace
 - Logically a mapping from pathname to chunk list
 - Allows concurrent file creation in same directory
 - Read/write locks prevent conflicting operations
 - File deletion by renaming to a hidden name; removed during regular scan
- Operation log
 - Historical record of metadata changes
 - Kept on multiple remote machines
 - Checkpoint created when log exceeds threshold
 - When checkpointing, switch to new log and create checkpoint in separate thread
 - Recovery made from most recent checkpoint and subsequent log
- Snapshot
 - Revokes leases on chunks in file/directory
 - Log operation
 - Duplicate metadata (not the chunks!) for the source
 - On first client write to chunk:
 - ❖ Required for client to gain access to chunk
 - ❖ Reference count > 1 indicates a duplicated chunk
 - ❖ Create a new chunk and update chunk list for duplicate

Chunk/replica management

- Placement
 - On chunkservers with below-average disk space utilization
 - Limit number of “recent” creations on a chunkserver (since access traffic will follow)
 - Spread replicas across racks (for reliability)
- Reclamation
 - Chunk become garbage when file of which they are a part is deleted
 - Lazy strategy (garbage collection) is used since no attempt is made to reclaim chunks at time of deletion
 - In periodic “HeartBeat” message chunkserver reports to the master a subset of its current chunks
 - Master identifies which reported chunks are no longer accessible (i.e., are garbage)
 - Chunkserver reclaims garbage chunks
- Stale replica detection
 - Master assigns a version number to each chunk/replica
 - Version number incremented each time a lease is granted
 - Replicas on failed chunkservers will not have the current version number
 - Stale replicas removed as part of garbage collection

Performance

Cluster	A	B
Chuckservers	342	227
Available disk space	72 TB	180 TB
Used disk space	55 TB	155 TB
Number of Files	735 k	737 k
Number of Dead files	22 k	232 k
Number of Chunks	992 k	1550 k
Metadata at chunkservers	13 GB	21 GB
Metadata at master	48 MB	60 MB



Cluster	A	B
Read rate (last minute)	583 MB/s	380 MB/s
Read rate (last hour)	562 MB/s	384 MB/s
Read rate (since restart)	589 MB/s	49 MB/s
Write rate (last minute)	1 MB/s	101 MB/s
Write rate (last hour)	2 MB/s	117 MB/s
Write rate (since restart)	25 MB/s	13 MB/s
Master ops (last minute)	325 Ops/s	533 Ops/s
Master ops (last hour)	381 Ops/s	518 Ops/s
Master ops (since restart)	202 Ops/s	347 Ops/s

HDFS

GOALS OF HDFS

- Very Large Distributed File System
 - 10K nodes, 100 million files, 10PB
- Assumes Commodity Hardware
 - Files are replicated to handle hardware failure
 - Detect failures and recover from them
- Optimized for Batch Processing
 - Data locations exposed so that computations can move to where data resides
 - Provides very high aggregate bandwidth

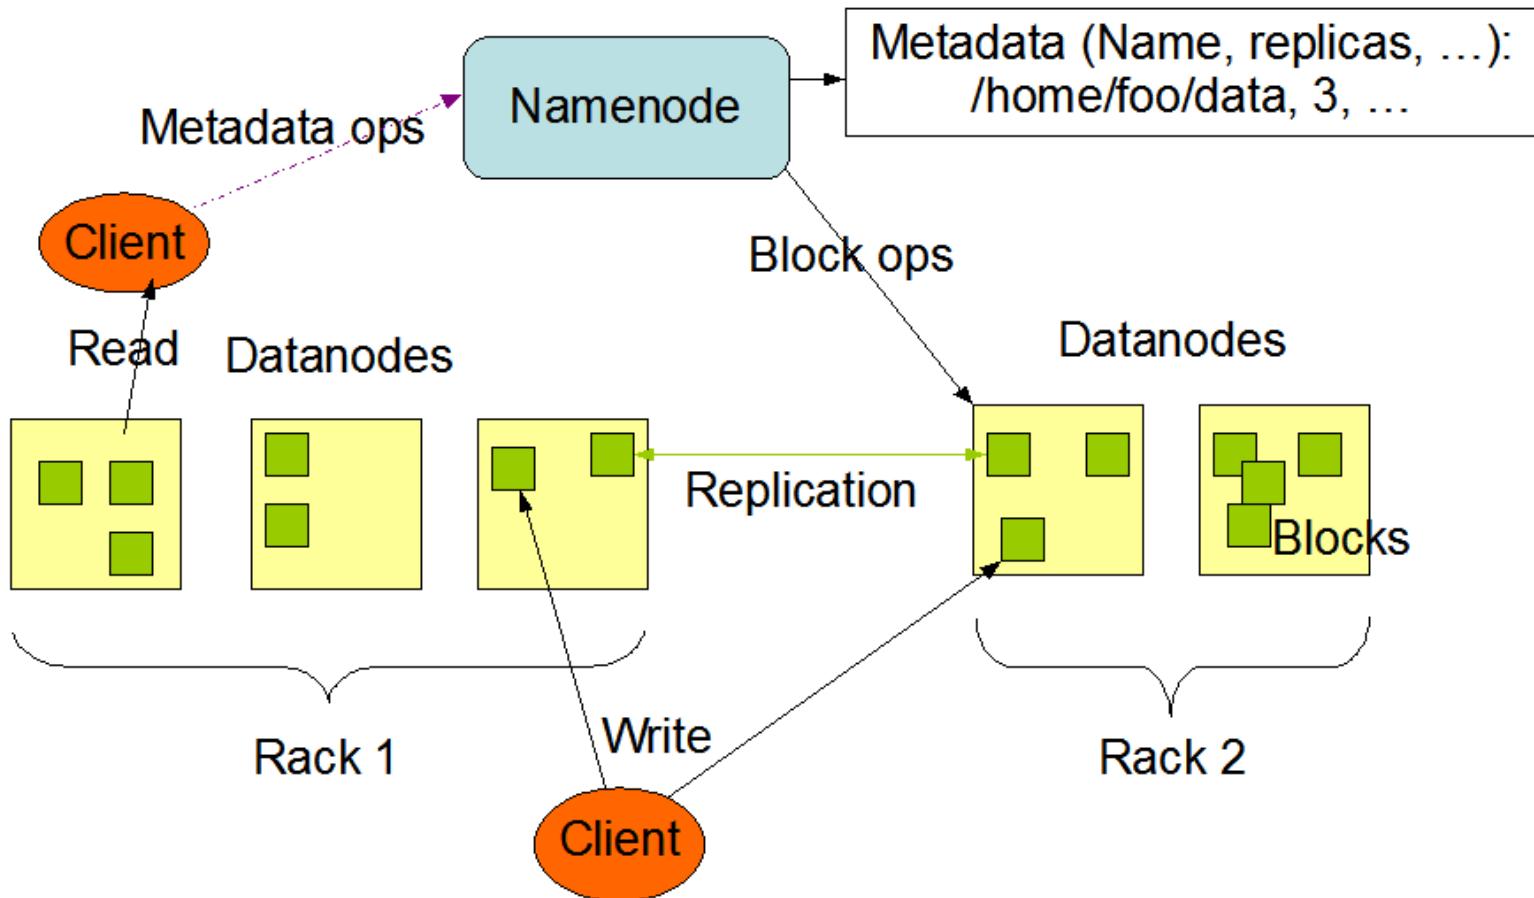


HADOOP DISTRIBUTED FILE SYSTEM

- ◉ Single Namespace for entire cluster
- ◉ Data Coherency
 - Write-once-read-many access model
 - Client can only append to existing files
- ◉ Files are broken up into blocks
 - Typically 64MB block size
 - Each block replicated on multiple DataNodes
- ◉ Intelligent Client
 - Client can find location of blocks
 - Client accesses data directly from DataNode

HDFS ARCHITECTURE

HDFS Architecture



FUNCTIONS OF A NAMENODE

- Manages File System Namespace
 - Maps a file name to a set of blocks
 - Maps a block to the DataNodes where it resides
- Cluster Configuration Management
- Replication Engine for Blocks

NAMENODE METADATA

- Metadata in Memory

- The entire metadata is in main memory
- No demand paging of metadata

- Types of metadata

- List of files
- List of Blocks for each file
- List of DataNodes for each block
- File attributes, e.g. creation time, replication factor

- A Transaction Log

- Records file creations, file deletions etc

DATANODE

- ◉ A Block Server

- Stores data in the local file system (e.g. ext3)
- Stores metadata of a block (e.g. CRC)
- Serves data and metadata to Clients

- ◉ Block Report

- Periodically sends a report of all existing blocks to the NameNode

- ◉ Facilitates Pipelining of Data

- Forwards data to other specified DataNodes

BLOCK PLACEMENT

- ◉ Current Strategy
 - One replica on local node
 - Second replica on a remote rack
 - Third replica on same remote rack
 - Additional replicas are randomly placed
- ◉ Clients read from nearest replicas
- ◉ Would like to make this policy pluggable

HEARTBEATS

- DataNodes send heartbeat to the NameNode
 - Once every 3 seconds
- NameNode uses heartbeats to detect DataNode failure

REPLICATION ENGINE

- NameNode detects DataNode failures
 - Chooses new DataNodes for new replicas
 - Balances disk usage
 - Balances communication traffic to DataNodes

DATA CORRECTNESS

- ◉ Use Checksums to validate data
 - Use CRC32
- ◉ File Creation
 - Client computes checksum per 512 bytes
 - DataNode stores the checksum
- ◉ File access
 - Client retrieves the data and checksum from DataNode
 - If Validation fails, Client tries other replicas

NAMENODE FAILURE

- ◉ A single point of failure
- ◉ Transaction Log stored in multiple directories
 - A directory on the local file system
 - A directory on a remote file system (NFS/CIFS)
- ◉ Need to develop a real HA solution

SECONDARY NAMENODE

- Copies FsImage and Transaction Log from Namenode to a temporary directory
- Merges FsImage and Transaction Log into a new FsImage in temporary directory
- Uploads new FsImage to the NameNode
 - Transaction Log on NameNode is purged

USER INTERFACE

- Commads for HDFS User:

- hadoop dfs -mkdir /foodir
- hadoop dfs -cat /foodir/myfile.txt
- hadoop dfs -rm /foodir/myfile.txt

- Commands for HDFS Administrator

- hadoop fsadmin -report
- hadoop fsadmin -decommision datanodename

- Web Interface

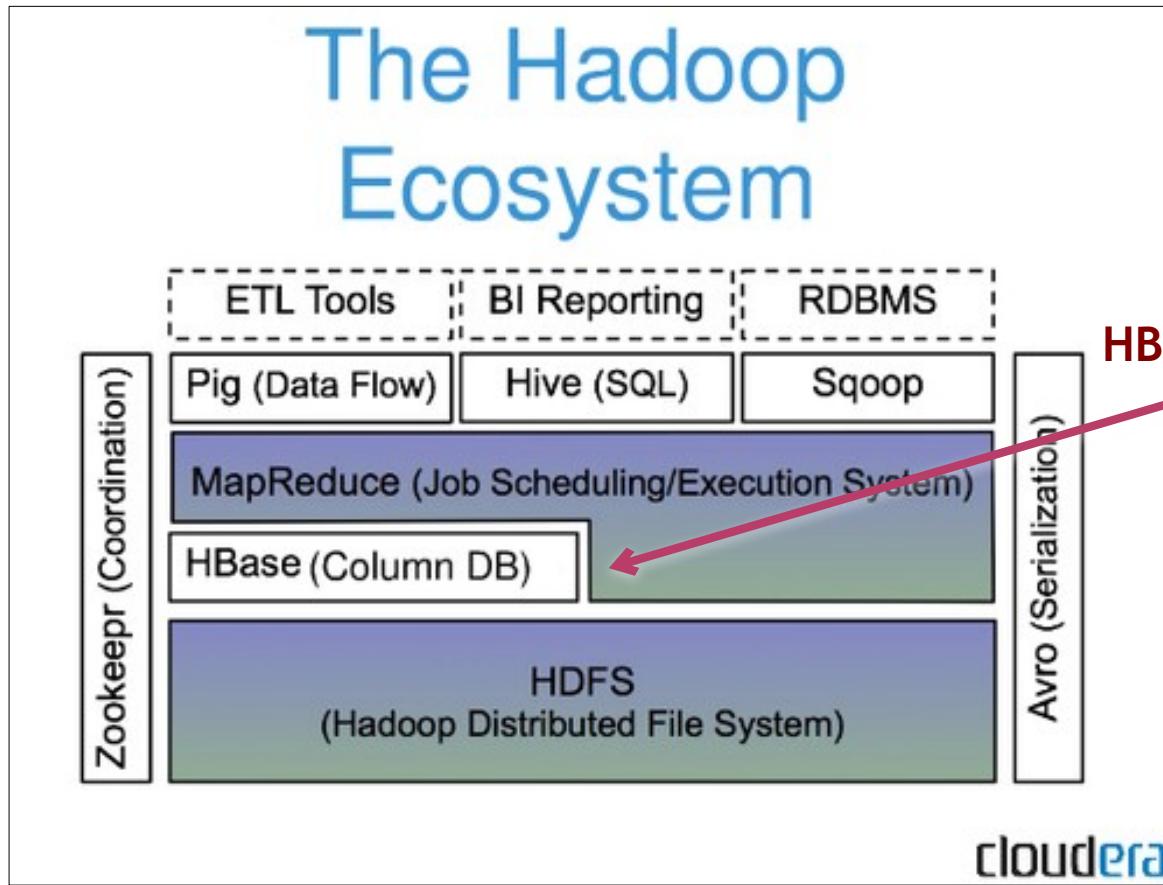
- <http://host:port/dfshealth.jsp>



HBASE: OVERVIEW

- **HBase is a distributed column-oriented data store built on top of HDFS**
- **HBase is an Apache open source project whose goal is to provide storage for the Hadoop Distributed Computing**
- **Data is logically organized into tables, rows and columns**

HBASE: PART OF HADOOP'S ECOSYSTEM



HBase is built on top of HDFS



HBase files are internally stored in HDFS

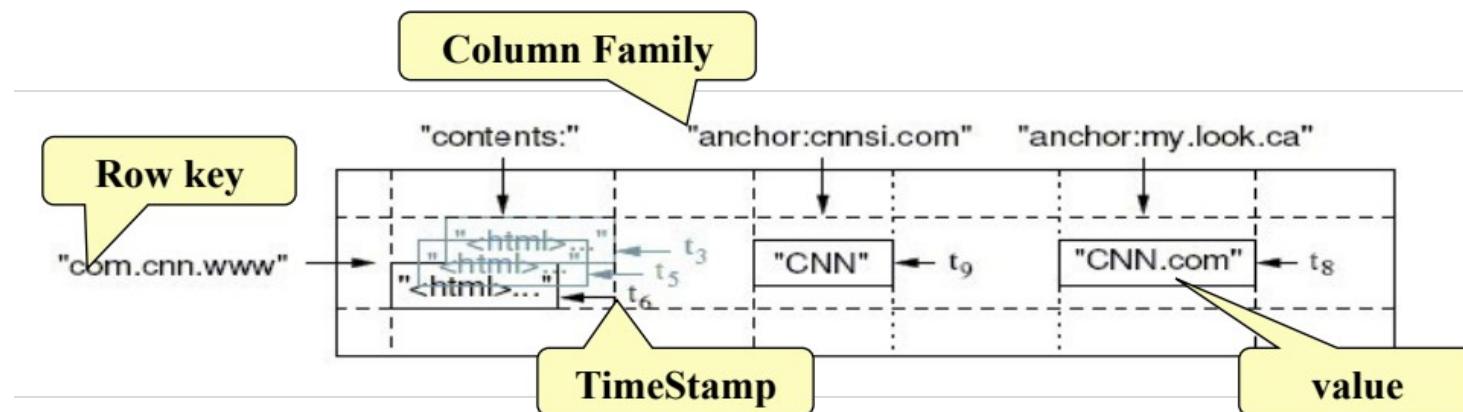
HBASE VS. HDFS (CONT'D)

- ◉ ***HBase*** is designed to efficiently address the above points
 - Fast record lookup
 - Support for record-level insertion
 - Support for updates (not in place)
- ◉ HBase updates are done by creating new versions of values

HBASE DATA MODEL

HBASE DATA MODEL

- ⦿ HBase is based on Google's Bigtable model
 - Key-Value pairs



HBASE LOGICAL VIEW

Implicit PRIMARY KEY in
RDBMS terms

Data is all byte [] in HBase

Different types of
data separated into
different
“column families”

Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer' @ts=2010, 'Hadoop': 'PMC' @ts=2011, 'Hive': 'Contributor' }

Different rows may have different sets
of columns (table is sparse)

A single cell might have different
values at different timestamps

Useful for *-To-Many mappings

HBASE: KEYS AND COLUMN FAMILIES

Each record is divided into Column Families

Each row has a Key

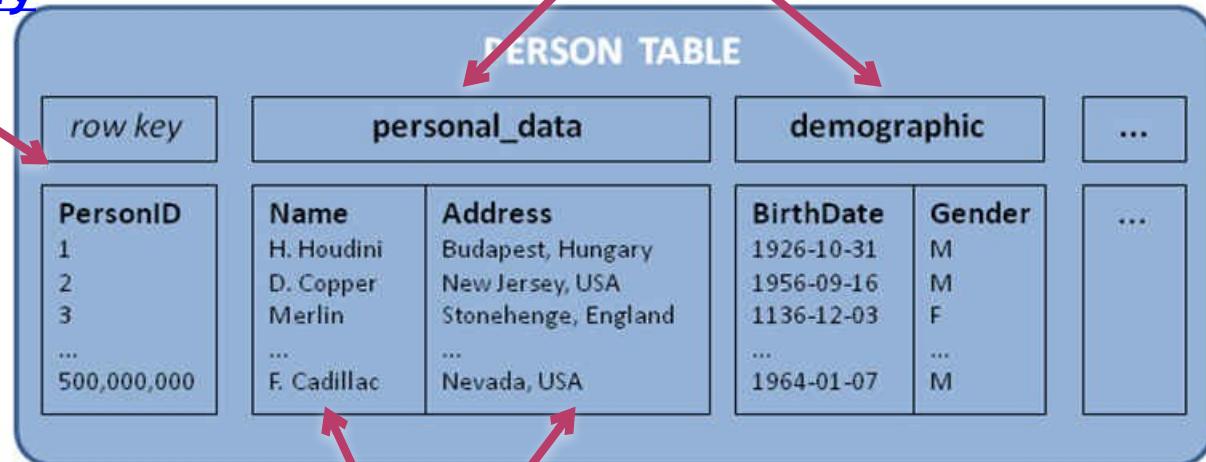


Figure 2. Census Data in Column Families

Each column family consists of one or more Columns

◎ Key

- Byte array
- Serves as the primary key for the table
- Indexed for fast lookup

◎ Column Family

- Has a name (string)
- Contains one or more related columns

◎ Column

- Belongs to one column family
- Included inside the row
 - *familyName:columnName*

Row key	Time Stamp	Column “content s:”	Column “anchor:”	
“com.apache.www”	t12	“<html> ...”		
	t11	“<html> ...”	Column named “apache.com”	
	t10		“anchor:apache.com”	“APACHE”
	t15		“anchor:cnnsci.com”	“CNN”
	t13		“anchor:my.look.ca”	“CNN.com”
	t6	“<html> ...”		
	t5	“<html> ...”		
	t3	“<html> ...”		

- **Version Number**
 - Unique within each key
 - By default → System's timestamp
 - Data type is Long
- **Value (Cell)**
 - Byte array

Version number for each row

Row key	Time Stamp	Column “content s:”	Column “anchor:”
“com.apac he.ww w”	t12	“<html> ...”	
	t11	“<html> ...”	
	t10		“anchor:apache .com” “APACH E”
	t15		“anchor:cnnsi.co m” “CNN”
	t13		“anchor:my.look. ca” “CNN.co m”
	t6	“<html> ...”	
“com.cnn.w ww”	t5	“<html> ...”	
	t3	“<html> ...”	

NOTES ON DATA MODEL

- HBase schema consists of several *Tables*
- Each table consists of a set of *Column Families*
 - Columns are not part of the schema
- HBase has *Dynamic Columns*
 - Because column names are encoded inside the cells
 - Different cells can have different columns

“Roles” column family has different columns in different cells



Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tlipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer'@ts=2010, 'Hadoop': 'PMC'@ts=2011, 'Hive': 'Contributor' }

NOTES ON DATA MODEL (CONT'D)

- The *version number* can be user-supplied
 - Even does not have to be inserted in increasing order
 - Version number are unique within each key
- Table can be very sparse
 - Many cells are empty
- *Keys* are indexed as the primary key

Has two columns
[cnnsi.com &
my.look.ca]

Row Key	Time Stamp	ColumnFamily contents	ColumnFamily anchor
"com.cnn.www"	t9		anchor:cnnsi.com = "CNN"
"com.cnn.www"	t8		anchor:my.look.ca = "CNN.com"
"com.cnn.www"	t6	contents:html = "<html>..."	
"com.cnn.www"	t5	contents:html = "<html>..."	
"com.cnn.www"	t3	contents:html = "<html>..."	

HBASE PHYSICAL MODEL

HBASE PHYSICAL MODEL

- Each column family is stored in a separate file (called *HTables*)
- Key & Version numbers are replicated with each column family
- Empty cells are not stored

HBase maintains a multi-level index on values:
<key, column family, column name, timestamp>

Table 5.3. ColumnFamily contents

Row Key	Time Stamp	ColumnFamily "contents:"
"com.cnn.www"	t6	contents:html = "<html>..."
"com.cnn.www"	t5	contents:html = "<html>..."
"com.cnn.www"	t3	contents:html = "<html>..."

Table 5.2. ColumnFamily anchor

Row Key	Time Stamp	Column Family anchor
"com.cnn.www"	t9	anchor:cnnsi.com = "CNN"
"com.cnn.www"	t8	anchor:my.look.ca = "CNN.com"

EXAMPLE

Row key	Data
cutting	info: { 'height': '9ft', 'state': 'CA' } roles: { 'ASF': 'Director', 'Hadoop': 'Founder' }
tlipcon	info: { 'height': '5ft7', 'state': 'CA' } roles: { 'Hadoop': 'Committer'@ts=2010, 'Hadoop': 'PMC'@ts=2011, 'Hive': 'Contributor' } ↑↑

info Column Family

Row key	Column key	Timestamp	Cell value
cutting	info:height	1273516197868	9ft
cutting	info:state	1043871824184	CA
tlipcon	info:height	1273878447049	5ft7
tlipcon	info:state	1273616297446	CA

roles Column Family

Row key	Column key	Timestamp	Cell value
cutting	roles:ASF	1273871823022	Director
cutting	roles:Hadoop	1183746289103	Founder
tlipcon	roles:Hadoop	1300062064923	PMC
tlipcon	roles:Hadoop	1293388212294	Committer
tlipcon	roles:Hive	1273616297446	Contributor

Sorted
on disk by
Row key, Col
key,
descending
timestamp

Milliseconds since unix epoch

cloudera

HBASE REGIONS

- Each HTable (column family) is partitioned horizontally into *regions*
 - Regions are counterpart to HDFS blocks

Table 5.3. ColumnFamily contents

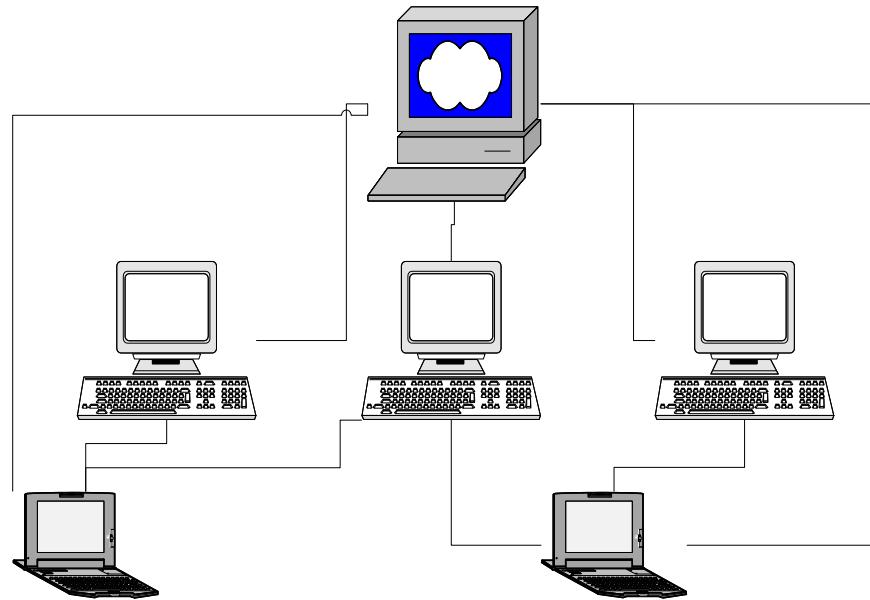
Row Key	Time Stamp	ColumnFamily "contents:"
"com.cnn.www"	t6	contents:html = "<html>..."
"com.cnn.www"	t5	contents:html = "<html>..."
"com.cnn.www"	t3	contents:html = "<html>..."

Each will be one region

HBASE ARCHITECTURE

THREE MAJOR COMPONENTS

- The HBaseMaster
 - One master
- The HRegionServer
 - Many region servers
- The HBase client



HBASE COMPONENTS

◎ Region

- A subset of a table's rows, like horizontal range partitioning
- Automatically done

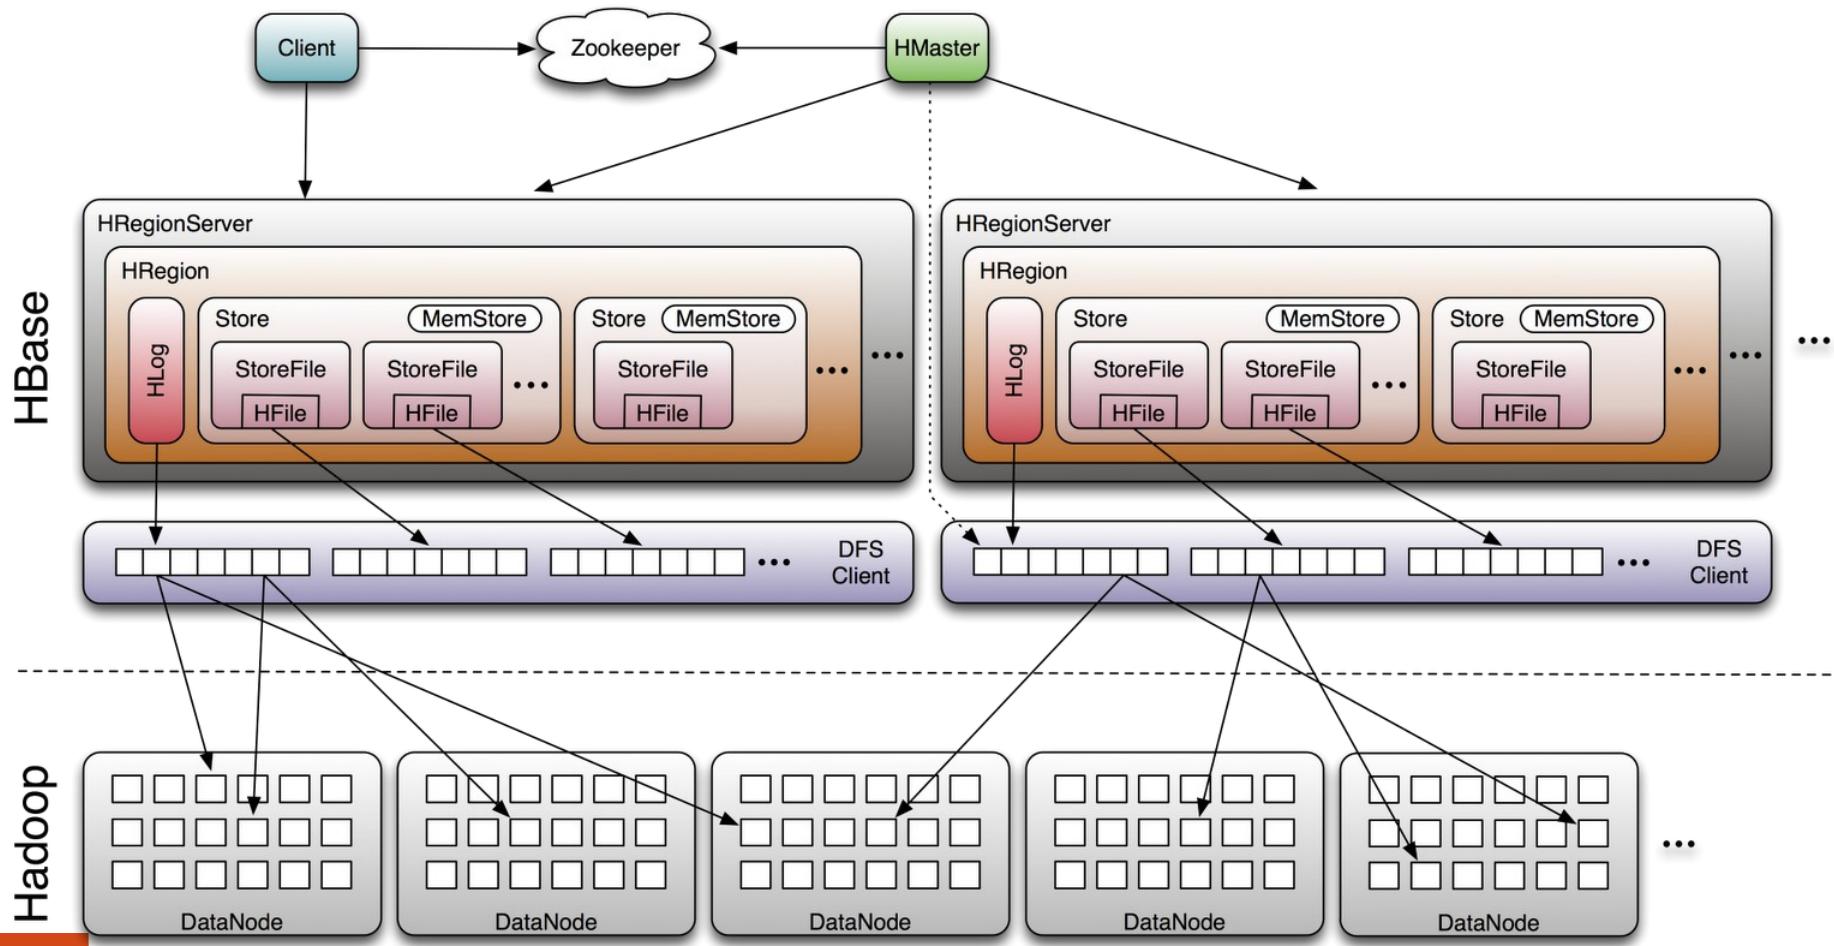
◎ RegionServer (many slaves)

- Manages data regions
- Serves data for reads and writes (*using a log*)

◎ Master

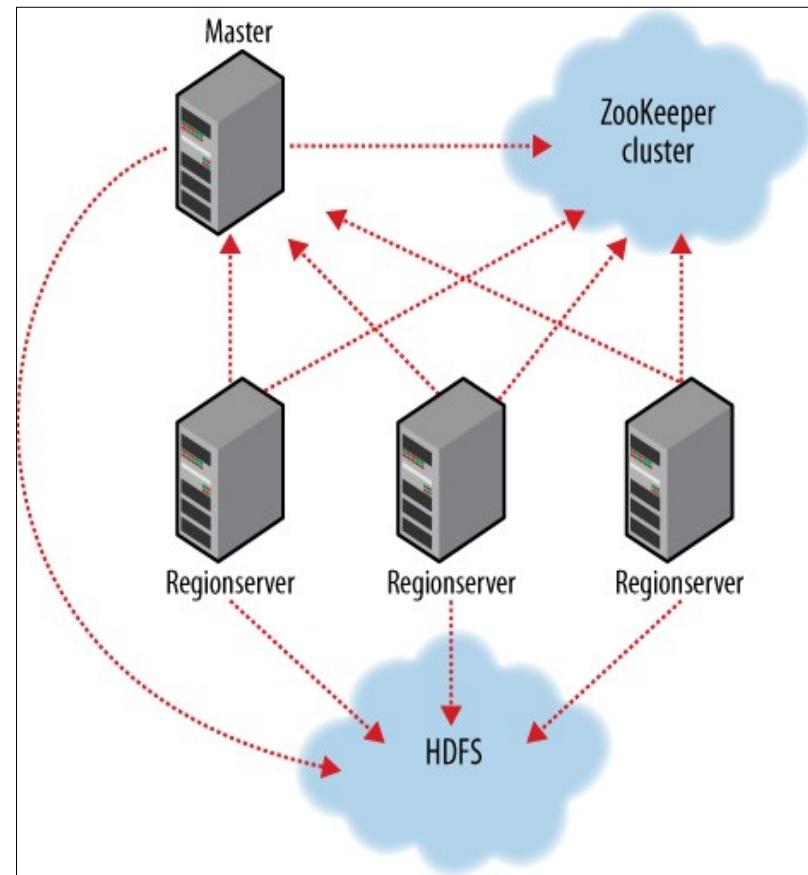
- Responsible for coordinating the slaves
- Assigns regions, detects failures
- Admin functions

BIG PICTURE



ZOOKEEPER

- HBase depends on ZooKeeper
- By default HBase manages the ZooKeeper instance
 - E.g., starts and stops ZooKeeper
- HMaster and HRegionServers register themselves with ZooKeeper



CREATING A TABLE

```
HBaseAdmin admin= new HBaseAdmin(config);
HColumnDescriptor []column;
column= new HColumnDescriptor[2];
column[0]=new HColumnDescriptor("columnFamily1:");
column[1]=new HColumnDescriptor("columnFamily2:");
HTableDescriptor desc= new
    HTableDescriptor(Bytes.toBytes("MyTable"));
desc.addFamily(column[0]);
desc.addFamily(column[1]);
admin.createTable(desc);
```

OPERATIONS ON REGIONS:

GET()

- Given a key → return corresponding record
- For each value return the highest version

```
Get get = new Get(Bytes.toBytes("row1"));
Result r = htable.get(get);
5.8.1.2. Default Get Example r.getValue(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns current version of value
```

- Can control the number of versions you want

```
Get get = new Get(Bytes.toBytes("row1"));
get.setMaxVersions(3); // will return last 3 versions of row
Result r = htable.get(get);
byte[] b = r.getValue(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns current version of value
List<KeyValue> kv = r.getColumn(Bytes.toBytes("cf"), Bytes.toBytes("attr")); // returns all versions of
```

OPERATIONS ON REGIONS: SCAN()

```
HTable htable = ...      // instantiate HTable

Scan scan = new Scan();
scan.addColumn(Bytes.toBytes("cf"),Bytes.toBytes("attr"));
scan.setStartRow( Bytes.toBytes("row"));           // start key is inclusive
scan.setStopRow( Bytes.toBytes("row" + (char)0)); // stop key is exclusive
ResultScanner rs = htable.getScanner(scan);
try {
    for (Result r = rs.next(); r != null; r = rs.next()) {
        // process result...
    } finally {
        rs.close(); // always close the ResultScanner!
    }
}
```

GET()

Select value from table where
key='com.apache.www' AND
label='anchor:apache.com'

Row key	Time Stamp	Column "anchor:"	
"com.apache.www"	t12		
	t11		
	t10	"anchor:apache.com"	"APACHE"
	t9	"anchor:cnnsi.com"	"CNN"
	t8	"anchor:my.look.ca"	"CNN.com"
	t6		
	t5		
	t3		

SCAN()

Select value from table
where
anchor='cnnsi.com'

Row key	Time Stamp	Column "anchor:"	
"com.apache.www"	t12		
	t11		
	t10	"anchor:apache.com"	"APACHE"
"com.cnn.www"	t9	"anchor:cnnsi.com"	"CNN"
	t8	"anchor:my.look.ca"	"CNN.com"
	t6		
	t5		
	t3		

OPERATIONS ON REGIONS:

PUT()

- Insert a new record (with a new key), Or
- Insert a record for an existing key

Implicit version number
(timestamp)

```
Put put = new Put(Bytes.toBytes(row));
put.add(Bytes.toBytes("cf"), Bytes.toBytes("attr1"), Bytes.toBytes( data));
htable.put(put);
```

Explicit version number

```
Put put = new Put( Bytes.toBytes(row));
long explicitTimeInMs = 555; // just an example
put.add(Bytes.toBytes("cf"), Bytes.toBytes("attr1"), explicitTimeInMs, Bytes.toBytes(data));
htable.put(put);
```

OPERATIONS ON REGIONS:

DELETE()

- Marking table cells as deleted
- **Multiple levels**
 - Can mark an entire column family as deleted
 - Can make all column families of a given row as deleted

- All operations are logged by the RegionServers
- The log is flushed periodically

HBASE: JOINS

- ◉ HBase does not support joins
- ◉ Can be done in the application layer
 - Using scan() and get() operations

ALTERING A TABLE

```
Configuration config = HBaseConfiguration.create();
HBaseAdmin admin = new HBaseAdmin(config);
String table = "myTable";
admin.disableTable(table);           ← Disable the table before changing the schema

HColumnDescriptor cf1 = ....;
admin.addColumn(table, cf1);          // adding new ColumnFamily
HColumnDescriptor cf2 = ....;
admin.modifyColumn(table, cf2);        // modifying existing ColumnFamily

admin.enableTable(table);
```

6.1. Schema Creation

Distributed Data Storage and Big Data

5.1 Distributed File Systems

- 5.1.1 Google File System (GFS):
Architecture, Operations, features and use cases
- 5.1.2 Hadoop Distributed File System (HDFS):
Architecture, Operations, features and use cases

5.2 NoSQL Databases:

- 5.2.1 Cassandra: Architecture, design principles,
Data Model and Query Language, features and use cases
- 5.2.2 MongoDB: Architecture, design principles,
Data Model and Query Language, features and use cases

5.3 Big Data Processing Frameworks

- 5.3.1 Hadoop: Key features and use cases
- 5.3.2 Spark : Key Features and use cases